

Model variability and Bias-Variance Trade-off in Decision Tree regressor

January 21, 2021

```
[1]: #import the libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

```
[2]: #import the datasets
df=pd.read_csv('insurance.csv')
```

```
[3]: df.head()
```

```
[3]:
```

	age	sex	bmi	children	smoker	region	expenses
0	19	female	27.9	0	yes	southwest	16884.92
1	18	male	33.8	1	no	southeast	1725.55
2	28	male	33.0	3	no	southeast	4449.46
3	33	male	22.7	0	no	northwest	21984.47
4	32	male	28.9	0	no	northwest	3866.86

```
[4]: df.tail()
```

```
[4]:
```

	age	sex	bmi	children	smoker	region	expenses
1333	50	male	31.0	3	no	northwest	10600.55
1334	18	female	31.9	0	no	northeast	2205.98
1335	18	female	36.9	0	no	southeast	1629.83
1336	21	female	25.8	0	no	southwest	2007.95
1337	61	female	29.1	0	yes	northwest	29141.36

```
[5]: df.describe(include='all')
```

```
[5]:
```

	age	sex	bmi	children	smoker	region	\
count	1338.000000	1338	1338.000000	1338.000000	1338	1338	
unique	NaN	2	NaN	NaN	2	4	
top	NaN	male	NaN	NaN	no	southeast	
freq	NaN	676	NaN	NaN	1064	364	
mean	39.207025	NaN	30.665471	1.094918	NaN	NaN	

std	14.049960	NaN	6.098382	1.205493	NaN	NaN
min	18.000000	NaN	16.000000	0.000000	NaN	NaN
25%	27.000000	NaN	26.300000	0.000000	NaN	NaN
50%	39.000000	NaN	30.400000	1.000000	NaN	NaN
75%	51.000000	NaN	34.700000	2.000000	NaN	NaN
max	64.000000	NaN	53.100000	5.000000	NaN	NaN

	expenses
count	1338.000000
unique	NaN
top	NaN
freq	NaN
mean	13270.422414
std	12110.011240
min	1121.870000
25%	4740.287500
50%	9382.030000
75%	16639.915000
max	63770.430000

```
[6]: #check for the null values
df.isnull().sum()
```

```
[6]: age      0
sex        0
bmi        0
children   0
smoker     0
region     0
expenses   0
dtype: int64
```

```
[7]: #visualize the null values
sns.heatmap(df.isnull()==True, cbar=False, yticklabels=False)
#It shows that there is no null values in the datasets
```

```
[7]: <matplotlib.axes._subplots.AxesSubplot at 0x2509871af70>
```



```
[8]: df.shape
```

```
[8]: (1338, 7)
```

```
[9]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   expenses    1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

```
[10]: # x and y variables
x=df.drop('expenses', axis=1)
```

```
[11]: y=df['expenses'].values
```

```
[12]: print(y)
```

```
[16884.92  1725.55  4449.46 ... 1629.83  2007.95 29141.36]
```

```
[13]: #Now convert the catagorical variables to numeric variables
```

```
cols=['sex', 'smoker', 'region']
```

```
x= pd.get_dummies(data=x, columns=cols, drop_first=True)
```

```
[14]: x.head()
```

```
[14]:
```

	age	bmi	children	sex_male	smoker_yes	region_northwest	\
0	19	27.9	0	0	1	0	
1	18	33.8	1	1	0	0	
2	28	33.0	3	1	0	0	
3	33	22.7	0	1	0	1	
4	32	28.9	0	1	0	1	

	region_southeast	region_southwest
0	0	1
1	1	0
2	1	0
3	0	0
4	0	0

```
[15]: #scale the features age and bmi to the same scale as of the other features
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
[16]: scaler=MinMaxScaler()
```

```
scale_cols=['age', 'bmi']
```

```
x[scale_cols]=scaler.fit_transform(x[scale_cols])
```

```
[17]: x.head()
```

```
[17]:
```

	age	bmi	children	sex_male	smoker_yes	region_northwest	\
0	0.021739	0.320755	0	0	1	0	
1	0.000000	0.479784	1	1	0	0	
2	0.217391	0.458221	3	1	0	0	
3	0.326087	0.180593	0	1	0	1	
4	0.304348	0.347709	0	1	0	1	

	region_southeast	region_southwest
0	0	1
1	1	0
2	1	0
3	0	0
4	0	0

```
[18]: from sklearn.tree import DecisionTreeRegressor
```

```
[19]: from sklearn.model_selection import train_test_split

[20]: x_train, x_test, y_train, y_test=train_test_split(x, y, test_size=0.25,
↳random_state=0)

[21]: tree_reg=DecisionTreeRegressor(max_depth=2, criterion='mse')
tree_reg.fit(x_train, y_train)

[21]: DecisionTreeRegressor(max_depth=2)

[22]: y_pred_reg=tree_reg.predict(x_test)

[23]: y_pred_reg.T

[23]: array([12428.30292135, 12428.30292135, 41512.0223301 , 12428.30292135,
12428.30292135, 5416.65981982, 5416.65981982, 12428.30292135,
5416.65981982, 5416.65981982, 5416.65981982, 12428.30292135,
12428.30292135, 5416.65981982, 21502.9989 , 12428.30292135,
12428.30292135, 5416.65981982, 5416.65981982, 41512.0223301 ,
21502.9989 , 12428.30292135, 12428.30292135, 21502.9989 ,
5416.65981982, 5416.65981982, 5416.65981982, 5416.65981982,
5416.65981982, 12428.30292135, 5416.65981982, 41512.0223301 ,
12428.30292135, 12428.30292135, 21502.9989 , 5416.65981982,
12428.30292135, 41512.0223301 , 41512.0223301 , 5416.65981982,
5416.65981982, 5416.65981982, 21502.9989 , 41512.0223301 ,
41512.0223301 , 5416.65981982, 12428.30292135, 5416.65981982,
5416.65981982, 12428.30292135, 5416.65981982, 5416.65981982,
21502.9989 , 41512.0223301 , 12428.30292135, 5416.65981982,
5416.65981982, 12428.30292135, 12428.30292135, 12428.30292135,
5416.65981982, 41512.0223301 , 12428.30292135, 12428.30292135,
12428.30292135, 12428.30292135, 41512.0223301 , 41512.0223301 ,
5416.65981982, 5416.65981982, 12428.30292135, 12428.30292135,
21502.9989 , 12428.30292135, 12428.30292135, 12428.30292135,
12428.30292135, 12428.30292135, 21502.9989 , 41512.0223301 ,
12428.30292135, 41512.0223301 , 5416.65981982, 12428.30292135,
41512.0223301 , 21502.9989 , 5416.65981982, 5416.65981982,
12428.30292135, 41512.0223301 , 5416.65981982, 12428.30292135,
5416.65981982, 12428.30292135, 5416.65981982, 5416.65981982,
41512.0223301 , 41512.0223301 , 5416.65981982, 12428.30292135,
5416.65981982, 5416.65981982, 5416.65981982, 41512.0223301 ,
21502.9989 , 5416.65981982, 12428.30292135, 5416.65981982,
12428.30292135, 41512.0223301 , 12428.30292135, 5416.65981982,
12428.30292135, 41512.0223301 , 41512.0223301 , 5416.65981982,
5416.65981982, 12428.30292135, 12428.30292135, 12428.30292135,
41512.0223301 , 12428.30292135, 12428.30292135, 5416.65981982,
12428.30292135, 5416.65981982, 21502.9989 , 21502.9989 ,
41512.0223301 , 5416.65981982, 12428.30292135, 5416.65981982,
```

[illegible]

```
5416.65981982, 5416.65981982, 5416.65981982, 12428.30292135,
5416.65981982, 12428.30292135, 41512.0223301 , 21502.9989 ,
12428.30292135, 41512.0223301 , 12428.30292135, 5416.65981982,
12428.30292135, 5416.65981982, 21502.9989 ])
```

0.0.1 How to Check the Model Variability and Bias-Variance Trade-Off??

```
[24]: from sklearn.model_selection import cross_validate, KFold
```

```
[25]: K_Fold=KFold(n_splits=10, shuffle=True, random_state=42)
depth={}

for i in range(2,11):
    tree_cv=cross_validate(DecisionTreeRegressor(max_depth=i), x, y, cv=K_Fold,
    scoring=['r2'])
    depth['depth_' +str(i)]=tree_cv['test_r2']
depth
```

```
[25]: {'depth_2': array([0.82797876, 0.83627513, 0.82846619, 0.79627887, 0.88576425,
0.82437852, 0.78101358, 0.73007204, 0.80832654, 0.83971832]),
'depth_3': array([0.85628937, 0.87237849, 0.85760018, 0.81877829, 0.91590791,
0.85343578, 0.83150173, 0.75706664, 0.80167779, 0.84678683]),
'depth_4': array([0.85300297, 0.86905526, 0.86086231, 0.84318317, 0.91869788,
0.84800633, 0.82755151, 0.74791778, 0.81648643, 0.85371445]),
'depth_5': array([0.84993277, 0.8802852 , 0.87258964, 0.8492467 , 0.89978505,
0.85098472, 0.80967563, 0.74832063, 0.81547351, 0.85134251]),
'depth_6': array([0.84655847, 0.85905411, 0.88005639, 0.82556283, 0.89192687,
0.84481696, 0.79283807, 0.7439385 , 0.80581394, 0.84064556]),
'depth_7': array([0.84696145, 0.830672 , 0.87047463, 0.79272067, 0.88495914,
0.83069584, 0.77334081, 0.73076654, 0.77484306, 0.78097405]),
'depth_8': array([0.82215768, 0.83557225, 0.83119287, 0.75026821, 0.8515351 ,
0.79491397, 0.72079798, 0.71526931, 0.78600994, 0.73315571]),
'depth_9': array([0.7837532 , 0.80745989, 0.81908771, 0.64583951, 0.84800343,
0.74552177, 0.64208233, 0.73292473, 0.75064992, 0.67133338]),
'depth_10': array([0.77312684, 0.78511287, 0.78250649, 0.66597412, 0.84996297,
0.67961468, 0.55739369, 0.68449381, 0.70355455, 0.69324033])}
```

```
[26]: #Now convert it into dataframe
df=pd.DataFrame(depth)
```


```
[27]: df
```

```
[27]:   depth_2  depth_3  depth_4  depth_5  depth_6  depth_7  depth_8 \
0  0.827979  0.856289  0.853003  0.849933  0.846558  0.846961  0.822158
1  0.836275  0.872378  0.869055  0.880285  0.859054  0.830672  0.835572
2  0.828466  0.857600  0.860862  0.872590  0.880056  0.870475  0.831193
3  0.796279  0.818778  0.843183  0.849247  0.825563  0.792721  0.750268
```

4	0.885764	0.915908	0.918698	0.899785	0.891927	0.884959	0.851535
5	0.824379	0.853436	0.848006	0.850985	0.844817	0.830696	0.794914
6	0.781014	0.831502	0.827552	0.809676	0.792838	0.773341	0.720798
7	0.730072	0.757067	0.747918	0.748321	0.743938	0.730767	0.715269
8	0.808327	0.801678	0.816486	0.815474	0.805814	0.774843	0.786010
9	0.839718	0.846787	0.853714	0.851343	0.840646	0.780974	0.733156

	depth_9	depth_10
0	0.783753	0.773127
1	0.807460	0.785113
2	0.819088	0.782506
3	0.645840	0.665974
4	0.848003	0.849963
5	0.745522	0.679615
6	0.642082	0.557394
7	0.732925	0.684494
8	0.750650	0.703555
9	0.671333	0.693240

```
[28]: # Now choose the best depth through visualizing boxplots
plt.figure(figsize=(10,10))
sns.set_style('whitegrid')
sns.boxplot(data=df)
plt.xticks(fontsize=14)
plt.yticks(fontsize=14)
plt.ylabel("R2 score", fontsize=14)
plt.title("Decision Tree variability and Bias-Variance Trade-off")
plt.show()
```

[29] : # Reference:
machine-learning_Dr.B. (2020, March 18). Lesson 17 Decision Trees.
YouTube. <https://www.youtube.com/watch?v=KIuB9nsVKqY&t=709s>